Yale University
CPSC 516, Spring 2023
Assignment 5

Chang Feng (Felix) Zhou cz397

P.1.

(a)

From our work in class, we know that $\nabla^2 f \succeq mI$ implies that f is m-strongly convex. This yields the inequality

$$f(y) - [f(x) + \langle \nabla f(x), y - x\rangle] \geq \frac{m}{2}\|y - x\|_2^2$$

by the definition of strong convexity.

On the other hand, we know by the second order Taylor expansion about x that

$$\begin{aligned} f(y) &= f(x) + \langle \nabla f(x), y - x\rangle + \frac{1}{2}(y-x)^T \nabla^2 f(\xi)(y-x) && \xi \in [x, y] \\ f(y) - f(x) + \langle \nabla f(x), x - y\rangle &\leq \frac{M}{2}\|y - x\|_2^2. && MI \succeq \nabla^2 f \end{aligned}$$

This shows both inequalities.

(b)

Suppose $f(z^*) = y^*$ and consider the inequality from P.1.(a)

$$y^* \leq f(z) \leq f(x) + \langle \nabla f(x), z - x\rangle + \frac{M}{2}\|z - x\|_2^2.$$

Let us minimize the RHS with respect to z by taking the derivative and setting it to 0. We must have

$$\begin{aligned} \nabla f(x) + M[z - x] &= 0 \\ z &= x - \frac{1}{M}\nabla f(x). \end{aligned}$$

Substituting this particular value of z to the RHS above yields

$$\begin{aligned} &f(x) + \left\langle \nabla f(x), -\frac{1}{M}\nabla f(x)\right\rangle + \frac{1}{2M}\|\nabla f(x)\|_2^2 \\ &= f(x) - \frac{1}{2M}\|\nabla f(x)\|_2^2 \end{aligned}$$

which is one of the desired inequalities.

To see the other inequality, We note that we wish to prove

$$\frac{1}{2m}\|\nabla f(x)\|_2^2 \geq [f(x) - f(z^*)]$$

which is known as the *Polyak-Lojasiewicz (PL)* condition.

By the definition of strong convexity,

$$\begin{aligned}
f(x) - f(z^*) &\le \langle \nabla f(x), x - z^* \rangle - \frac{m}{2}\|x - z^*\|_2^2 \\
&= \langle \nabla f(x), x - z^* \rangle - \frac{m}{2}\|x - z^*\|_2^2 - \frac{1}{2m}\|\nabla f(x)\|_2^2 + \frac{1}{2m}\|\nabla f(x)\|_2^2 \\
&= -\frac{1}{2}\left\|\sqrt{m}(x - z^*) - \frac{1}{\sqrt{m}}\nabla f(x)\right\|_2^2 + \frac{1}{2m}\|\nabla f(x)\|_2^2 \\
&\le \frac{1}{2m}\|\nabla f(x)\|_2^2.
\end{aligned}$$

Having shown both inequalities, we conclude the proof.

(c)

In P.1.(b), we have shown that

$$f(z) \le f(x) - \frac{1}{2M}\|\nabla f(x)\|_2^2$$

for $z := x - \frac{1}{M}\nabla f(x)$.

But since we chose the step size α to minimize $f(x_{t+1})$, it must be at least as good as $\alpha = \frac{1}{M}$. Thus

$$f(x_{t+1}) \le f(x_t) - \frac{1}{M}\|\nabla f(x_t)\|_2^2$$

as desired.

(d)

We argue by induction, the base case

$$f(x_0) - y^* \le 1 \cdot [f(x_0) - y^*]$$

holds trivially. Suppose it holds up to some t and consider $f(x_{t+1}) - y^*$.

We have

$$\begin{aligned}
f(x_{t+1}) - y^* &= f(x_{t+1}) - f(x_t) + f(x_t) - y^* \\
&\le f(x_t) - y^* - \frac{1}{2M}\|\nabla f(x_t)\|_2^2 && \text{P.1.(c)} \\
&\le f(x_t) - y^* - \frac{m}{M}[f(x_t) - y^*] && \text{P.1.(b) LHS} \\
&= \left(1 - \frac{m}{M}\right)[f(x_t) - y^*] \\
&= \left(1 - \frac{m}{M}\right)^{t+1}[f(x_0) - y^*]. && \text{induction hypothesis}
\end{aligned}$$

By induction, we conclude the proof.

The number of iterations to reach ε error can be computed as follows

$$\begin{aligned}
\left(1-\frac{m}{M}\right)^t [f(x_0)-y^*] &\leq \exp(-mt/M)[f(x_0)-y^*] && 1-x\leq e^{-x}\\
&\leq \varepsilon \\
-\frac{mt}{M}+\log[f(x_0)-y^*] &\leq \log\varepsilon \\
t &\geq \frac{M}{m}\log\frac{f(x_0)-y^*}{\varepsilon}.
\end{aligned}$$

(e)

Suppose we are given A, b as input.

Consider the minimization problem

$$\begin{aligned}
&\min \|Ax-b\|_2^2\\
&x\in\mathbb{R}^n
\end{aligned}$$

The objective is the composition of an affine function and a convex, separable, and non-decreasing (in each coordinate) function, which is therefore convex.

We explicitly compute its first and second derivatives

$$\begin{aligned}
\frac{d}{dx}[Ax+b]^T[Ax+b] &= \frac{d}{dx}xA^2x+2x^TAb+b^Tb\\
&= 2A^2x+2Ab\\
\frac{d^2}{dx^2} &= 2A^2.
\end{aligned}$$

The objective is certainly twice differentiable, and since the eigenvalues of A^2 are just the eigenvalues of A squared,

$$\lambda_1(A)^2 \leq \boldsymbol{\nabla} f^2 \leq \lambda_n(A)^2.$$

By our work above, if we start with an initial solution $x_0 := 0$ and run gradient descent with step size $\alpha = \frac{1}{\lambda_n(A)^2}$, this yields a solution x such that $\|Ax-b\|_2^2 \leq \varepsilon$ after

$$T = O\left(\frac{\lambda_n(A)^2}{\lambda_1(A)^2}\log\frac{\|b\|_2^2}{\varepsilon}\right)$$

iterations. In each iteration, we need to compute the gradient and subtract it from the current iterate. The number of arithmetic operations is dominated by the gradient computation $A(Ax)$, which requires $O(n^2)$ operations if we compute Ax and then $A(Ax)$.

Thus the algorithm terminates after performing

$$O(n^2T) = O\left(n^2\kappa^2\log\frac{\|b\|_2^2}{\varepsilon}\right)$$

arithmetic operations.

P.2.