

Lecture 1

Algorithms via Convex Optimization


CPSC 368/516, Spring 2023

Nisheeth Vishnoi

Yale

Administrative Stuff

- Tuesday: 9.25 – 11.15
- Professor: Nisheeth VISHNOI
 - 10 Hillhouse, Room 227
 - Appointment by email
- TA: Anay Mehrotra
 - a.mehrotra@yale.edu
 - Office hours: Thursday 4-5 PM? (on Zoom)
- Course will be largely based on the book:
<https://convex-optimization.github.io/>
- CANVAS – everyone must register!



Content

- **Turing machines**
- **Part I – Convexity**
 - Basics of calculus, linear algebra, probability, ...
 - Convexity
 - Convex programming and efficiency
- **Part II – 1st-order methods for convex optimization (with applications)**
 - Gradient descent (flows/cuts)
 - Mirror descent and multiplicative weights method (matching)
- **Part III – Second-order/advanced methods (with applications)**
 - Newton's method
 - Interior point methods (linear programming, flows)
 - Ellipsoid methods (submodular functions, counting)

Content

- **Mathematical:** Significant experience in mathematical problem solving, writing proofs. Must solve homework problems, write them up (ideally in latex) and submit
- **Prerequisites:** Calculus, linear algebra, and probability, or permission of the instructor
- **What this course is not?**
 - *A first course in proofs/discrete mathematics*
 - *An introduction to machine learning*
- **End Goal:** Prepare you for **mathematical** research in theoretical computer science, optimization, and machine learning

Grading - Undergraduates

- **Problem sets** – 40% (~ 8 problem sets/4 graded)
- **Exam 1** – 30% (Week of March 6)
- **Exam 2** – 30% (Week of April 24)

Grading - Graduates


- **Problem sets – 30%** (~ 8 problem sets/4 graded)
- **Exam 1 – 30%** (Week of March 6)
- **Exam 2 – 30%** (Week of April 24)
- **Additional work – 10%**

Discrete problems in TCS/Optimization

- **Shortest path**

Input: Graph $G = (V, E)$, source s , and sink t


Output: Shortest “path” from s to t



- **s-t-Max Flow**

Input: Graph $G = (V, E)$, source s , and sink t


Output: Maximum “flow” from s to t such that at most **1** unit flow per edge



- **Bipartite Matching**

Input: Graph $G = (L, R, E)$


Output: Decide if G has a perfect matching



- **Count spanning trees**

Input: Graph $G = (V, E)$

Output: Count the number of spanning trees in G



Old and new approaches

- **Old Idea:**

- Formulate an optimization problem over **discrete variables**
- Use **combinatorial/discrete** optimization methods

- **New approach:**

- Formulate a (convex) formulation over **continuous domains**
- Use **continuous** methods (convex optimization)
- Prove correctness, establish precise running time guarantees

- **Why?**

- Big data – old algorithms may be **slow**
- Combination of this idea with tools such as linear solvers have led to **fastest known algorithms** for *nearly all* discrete optimization problems
- **Added benefits:**
 - Learn methods important in many areas (e.g., ML)

The s - t -maximum flow problem

s - t -maximum flow problem captures many **discrete optimization problems**, e.g., generalizes **bipartite matching**, **scheduling**, **routing**


Input: 1) Undirected graph $G = (V, E)$, $n := |V|$, $m := |E|$
2) Source and sink $s, t \in V$, $s \neq t$

Vertex-edge incidence matrix $B \in \mathbb{R}^{n \times m}$

$\forall i \in E$, direct $i := (u, v)$, B has a column $b_i := e_u - e_v$

Output: s - t -flow $x: E \rightarrow \mathbb{R}$ satisfies

- 1) **flow conservation**: for all $j \in V \setminus \{s, t\}$, $\langle e_j, Bx \rangle = 0$
- 2) **feasibility**: for all $i \in E$, $|x_i| \leq 1$ (**capacity 1**)



$$B = \begin{bmatrix} -1 & -1 & 0 & 0 & 0 \\ +1 & 0 & -1 & 0 & -1 \\ 0 & +1 & +1 & -1 & 0 \\ 0 & 0 & 0 & +1 & +1 \end{bmatrix}$$

B

Problem: Find a feasible s - t -flow x that maximizes the flow out of s : $|\langle e_s, Bx \rangle|$

Fact: There exists an integral s - t -maximum flow $x_i \in \{-1, 0, 1\}$

B is **totally unimodular** \Rightarrow every sq. submatrix A of B satisfies $\det(A) \in \{-1, 0, 1\}$

Many combinatorial algorithms: **Ford-Fulkerson**, **Edmonds-Karp**, **Dinic**, ...

For the s - t -maximum flow problem with **capacity** $U \in \{1, 2, \dots\}$:

[Goldberg and Rao, 1998]: An $\tilde{O}(m \min(n^{2/3}, m^{1/2}) \log U)$ time exact algorithm for s - t -maximum flow. E.g., when $m = O(n)$ and $U = O(1)$, running time is $O(m^{1.5})$

Convex programming (continuous) approach for maxflow

s - t -maximum flow **reduces** to: Given $F \in \mathbb{R}$ find an s - t flow x of value at least F
(F can be found in $O(\log m)$ steps using binary search)

Idea 1: s - t — F flow is the same as finding a **point** in

$$\underbrace{\{x \in \mathbb{R}^m : Bx = F(e_s - e_t)\}}_{(K_1) \text{ } x \text{ is } s\text{-}t\text{-flow } F} \cap \underbrace{\{x \in \mathbb{R}^m : |x_i| \leq 1, \forall i \in [m]\}}_{(K_2) \text{ } x \text{ satisfies "capacities"}}$$

K_1 and K_2 are **convex sets**—they are defined by **linear equalities/inequalities**


Idea 2: Formulate as convex program. E.g.,

- 1) Find $x \in K_1$ that minimizes “distance” to K_2
- 2) Find $x \in K_2$ that minimizes “distance” to K_1

Both are **convex programs** (i) K_1, K_2 are **convex**, (ii) **distance to convex sets is convex**

Idea 3 [Lee-Rao-Srivastava, 2013]: Consider **nonlinear convex program**

$$\begin{aligned} \min_{x \in \mathbb{R}^m} \quad & \text{dist}(x, K_2) \\ \text{s.t.}, \quad & x \in K_1, \end{aligned}$$



where $\text{dist}(x, K_2)$ is the (squared) **Euclidean distance** between x and K_2

How do we solve the above convex program?

First-order methods for minimizing convex fns

Roughly, family of iterative methods: each step moves in direction of **negative gradient**

Theorem: Given $\varepsilon > 0$, **convex** function $f: \mathbb{R}^m \rightarrow \mathbb{R}$, and access to gradients of f the following gradient descent methods make $O(T)$ calls to the gradients of f and output a point $x \in \mathbb{R}^m$ such that

$$f(x) \leq f(x^*) + \varepsilon \quad (x^* - \text{optimal point})$$

Where

- **Gradient descent** assumes that f is **L -Lipschitz continuous** and has $T = O(L\varepsilon^{-1})$
- **Mirror-descent** assumes that **norm of gradient** of f is $\leq G$ and has $T = O(G^2\varepsilon^{-2})$
- **Accelerated GD** assumes that f is **L -Lipschitz continuous** and has $T = O(\sqrt{L\varepsilon^{-1}})$

[Lee-Rao-Srivastava, 2013] use **accelerated GD** to give:

An $\tilde{O}(mn^{1/3}\varepsilon^{-1/3})$ time algorithm that for any $\varepsilon > 0$, $F \in \mathbb{R}$ outputs a s - t -flow of value $\geq (1 - \varepsilon)F$. E.g., when $m = O(n)$, runtime $O(m^{4/3})$ (beats Goldberg-Rao)

It can be converted to an **exact algorithm**, but requires $\varepsilon \approx O(1/F)$

For general capacity graphs, **F can be large** – so the running time is **slow**..


Problem: There is convex, L -Lipschitz cont. f for which any GD-method has $T = \Omega(\sqrt{L\varepsilon^{-1}})$

Can we develop algorithms whose runtime scales as $O(\log \varepsilon^{-1})$?

Linear programming approach to maxflow

s - t -maximum flow is also **special case** of linear programming:

- (i) objective is to **maximize F**
- (ii) subject to **linear equality/inequality constraints**



Linear program: Given matrix $A \in \mathbb{R}^{n \times m}$, constraint vector $b \in \mathbb{R}^n$, a cost vector $c \in \mathbb{R}^m$, solve:

$$\begin{aligned} \min_{x \in \mathbb{R}^m} & \langle c, x \rangle \\ \text{s.t.} & Ax = b \text{ and } x \geq 0 \end{aligned}$$

Combinatorial algorithms for s - t -maximum flow **rely** on

- **max-flow min-cut theorem**,
- **integrality** of s - t -maximum flow

Linear prog. duality generalizes the max-flow min-cut theorem; e.g., [Farkas, 1902]

Dual of the above program:

$$\max_{y \in \mathbb{R}^n} \langle b, y \rangle, \text{ s.t. } A^T y \geq c$$

Theorem: For any matrix $A \in \mathbb{R}^{n \times m}$, constraint vector $b \in \mathbb{R}^n$, a cost vector $c \in \mathbb{R}^m$, if both primal and dual programs are feasible, then their optimal values are equal

But general linear programs—among other properties—do **not** guarantee **integrality**!

How to solve linear prog. in time polynomial in the bit-complexity of A, b, c ?


Ellipsoid method: LP is in P

[Khachiyan, 1979] A “geometric” algorithm to check **feasibility** of linear programs

- Along with **binary search**, gives an algorithm to **solve** a linear program

Requires: **Separation oracle** for $K := \{x: Ax = b, x \geq 0\}$

- Input: A point $x \in \mathbb{R}^n$
- Output: YES if x is in K , otherwise
 - A **certificate**—hyperplane H —separating x and K



Input: An Ellipsoid E containing K

At each iteration, guess the **center of E** as a point in K

Then, **update E** based on the **response** of the separation oracle

Key points: At each iteration

- the **volume** of E **reduces** sufficiently
- solves one **linear system**

Theorem: A **poly(L)** iteration algorithm for solving linear programs, where L is the bit-complexity of (A, b, c) . In **each iteration**, the algorithm makes **one call** to the separation and takes additional **poly(L) time**

But for s - t -maximum flow it is **slower** than [Goldberg and Rao, 1998]


Interior point methods: Faster LP algorithms

[Karmarkar, 1984] A **faster** algorithm for linear programming than Ellipsoid method

Main idea: “Convert” LP to an **unconstrained convex prog.** using **barrier functions**

Barrier function (Informal): A **convex** fn that is finite in interior of set and increases to infinity as one approaches the boundary

Example: For $Ax \leq b$, $F(x) := -\sum_i \log(b_i - \langle A_i, x \rangle)$



[Renegar, 1988] Combined the barrier-approach with **Newton's method**—a **second order optimization method**—to improve the running time

Input: A barrier function $F(x)$, and second-order oracle of $F(x)$

Main step: Minimize $\eta \langle c, x \rangle + F(x)$, for fixed $\eta > 0$ (also change η over time)

Theorem: A $\tilde{O}(\sqrt{m} \cdot L)$ step algorithm for solving linear programs, where L is the bit-complexity of (A, b, c) . In each step, the algorithm solves an $m \times m$ linear system

For s - t -maximum flow:

Theorem: [Lee and Sidford, 2014] An $\tilde{O}(mn^{1/2} \cdot \log^2 U)$ time algorithm for s - t -maximum flow problem. E.g., for any $m > n$ it is **faster** than $\tilde{O}(m^{1.5})$

Recently [Chen, Kyng, Liu, Peng, Probst, Sachdeva 2022] running time to $\tilde{O}(m)$!

Ellipsoid method for convex programs

Problem: Given **convex set** $K \subseteq \mathbb{R}^m$ and **convex function** $f: \mathbb{R}^m \rightarrow \mathbb{R}$: $\min_{x \in K} f(x)$

Ellipsoid method can be used to solve the most general convex programs

Theorem: $\text{poly}((T_K + T_f) \cdot m \cdot \log \varepsilon^{-1})$ time algorithm that outputs $x \in K$, s.t.
$$f(x) \leq f(x^*) + \varepsilon,$$

where T_K and T_f are the running time of separation oracle for K and first-order of f

Implies efficient algorithms for comb. problems; e.g., via submodular minimization

A **submodular (set-)function** $f: 2^{[m]} \rightarrow \mathbb{R}$ satisfies: For sets $S \subseteq T \subseteq [m]$ and $i \in [m]$,
$$f(S \cup \{i\}) - f(S) \geq f(T \cup \{i\}) - f(T)$$

Problem: Given submodular function $f: 2^{[m]} \rightarrow \mathbb{R}$ find its minimizer: $\text{argmin}_{S \subseteq [m]} f(S)$

Applications:

- Originated in discrete optimization, e.g., minimum s - t -cut in graphs
- Machine learning: Arises in objectives for data summarization, influence maximization

Theorem: There is an algorithm that, given oracle access to a submodular function f , and $\varepsilon > 0$, outputs $S \subseteq [m]$ such that

$$f(S) \leq f(S^*) + \varepsilon,$$

where S^* is minimizer of f . The algorithm makes $\text{poly}(m, \log(\varepsilon^{-1}))$ queries to f

Applications: Max-entropy distributions

Convex programming for (approximately) counting discrete objects


Counting problem: Given $G = (V, E)$, compute the number of spanning trees of G .

Let \mathcal{T}_G be the set of all spanning trees of G

Let P_G be spanning tree polytope, i.e., the convex hull of all spanning trees in \mathcal{T}_G

Optimization problem: Given $G = (V, E)$ and $\theta \in P_G$, write θ as a convex combination of the vertices of P_G so that the probability distribution corresponding to the convex combination maximizes the **Shannon entropy**:

$$\begin{aligned} \min_p \quad & - \sum_{T \in \mathcal{T}_G} p_T \log p_T \\ \text{s.t.} \quad & \sum_{T \in \mathcal{T}_G} p_T v_T = \theta, \sum_{T \in \mathcal{T}_G} p_T = 1, p_T \geq 0 \quad \forall T \in \mathcal{T}_G \end{aligned} \quad (1)$$



Connection: If θ is the average of the vertices in P_G , then the value of Prog. (1) is $\log |\mathcal{T}_G|$

Prog. (1) is convex – however, it has **exponentially** many variables


- e.g., for the complete graph Prog. (1) has $|\mathcal{T}_G| = n^{n-2}$ variables

The dual of Prog. (1) has n variables and can be efficiently solved using the Ellipsoid method [Singh and Vishnoi, 2014] [Straszak and Vishnoi, 2019]


Turing Machines

What is a computer?

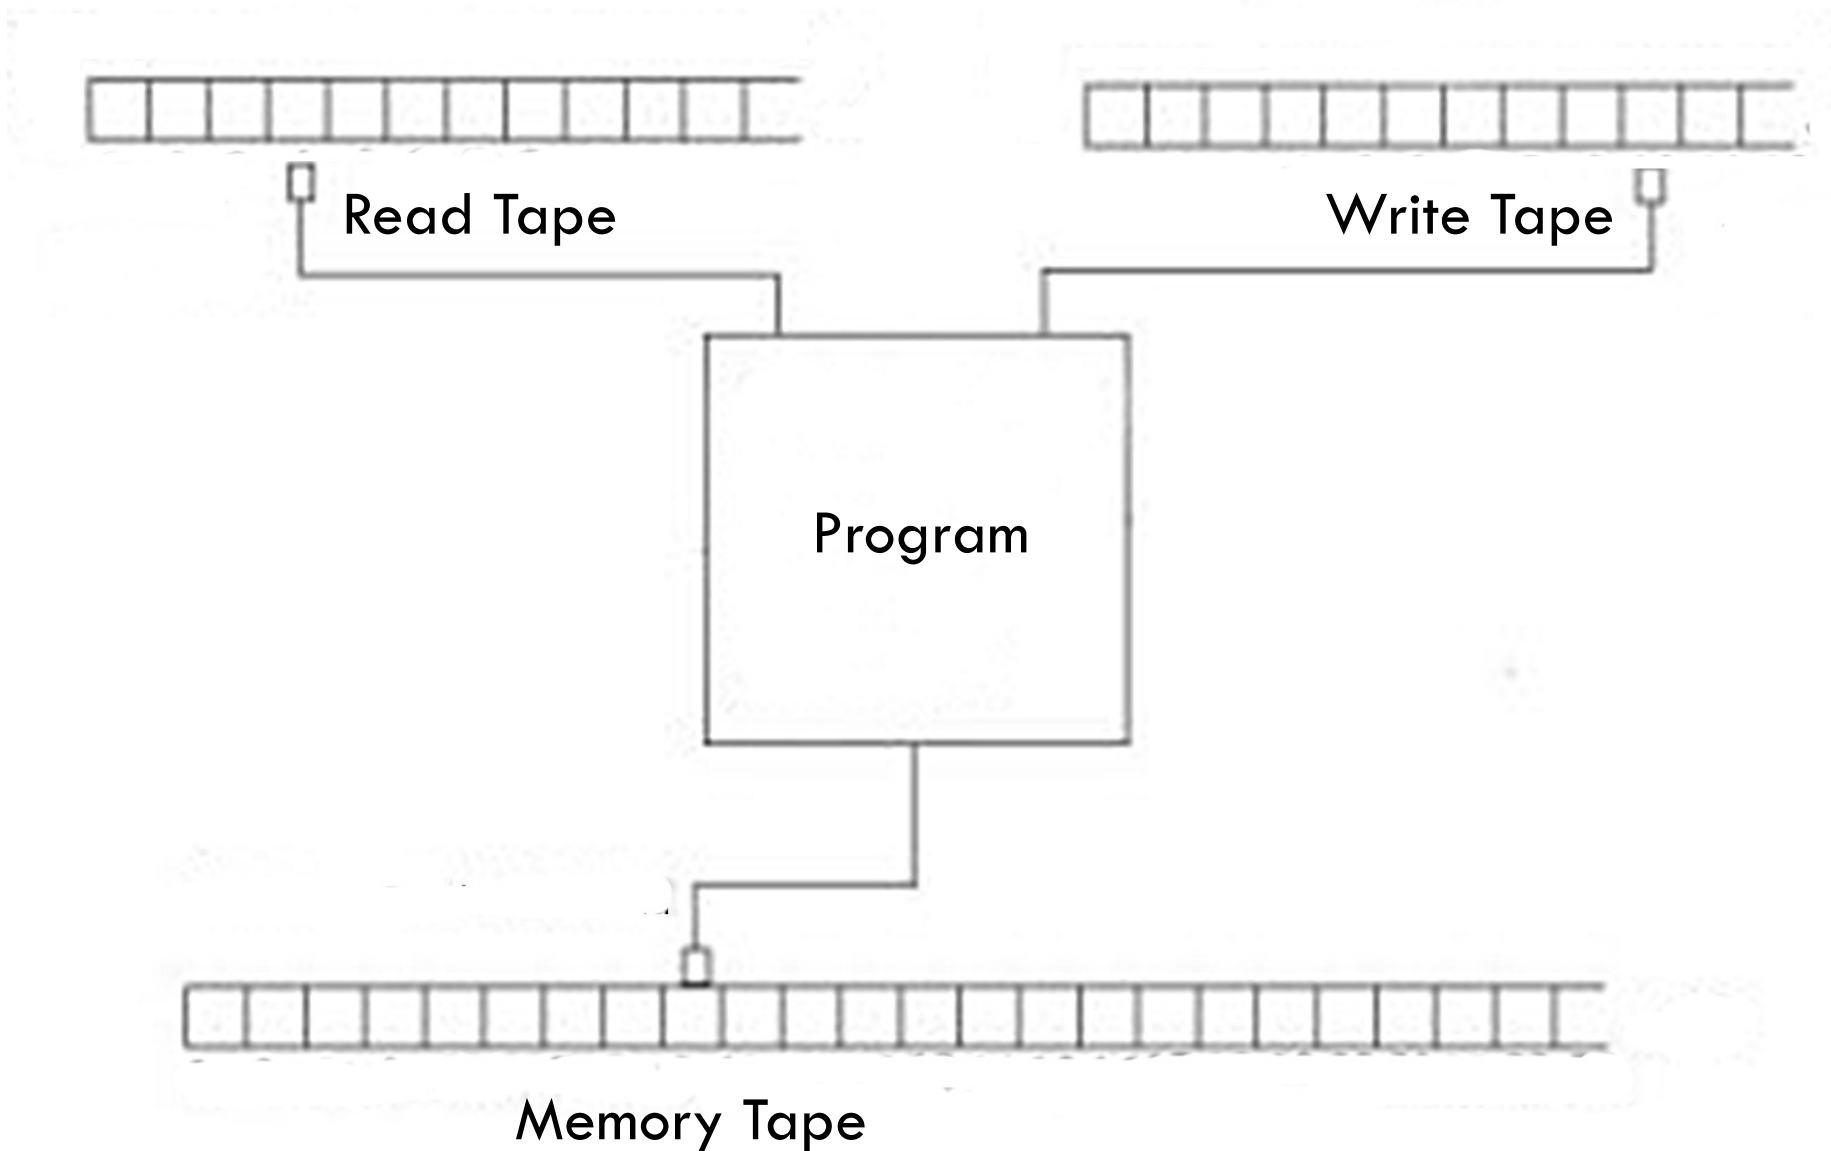
Program




Abstractly ..



More abstractly ..






Single tape seems enough ...



Program vs Finite Automata

```
1 def fib_tail(n):
2   def fib_tail_rec(a, b, n):
3     if n < 1:
4       return a
5     return fib_tail_rec(b, a + b, n - 1)
6   return fib_tail_rec(0, 1, n)
7
8
9
10 def fib_exponential(n):
11   if n == 0 or n == 1:
12     return n
13   else:
14     return fib_slow(n - 1) + fib_slow(n - 2)
15
```



**Abstract (special purpose)
computer**

Finite size program, larger and larger instances




Infinite Tape!

Infinite Tape

1	0	0	0	1	1	1	0			...
---	---	---	---	---	---	---	---	--	--	-----




Turing Machine



Head can **Read/Write/Move** Left/Right/Stay
 Once it reaches left-most cell, it can't go more left


Head has (finitely many) states



Exactly **one** Accept state
 and exactly **one** **Reject** state


Remaining states:
"computation in progress"

Tape Alphabet contains Input Alphabet



Example of starting configuration

May never reach an accept/reject state
May never HALT!



Formal Definition of a TM

A **Turing Machine** is a 7-tuple, $(Q, \Sigma, \Gamma, \delta, q_0, q_{\text{accept}}, q_{\text{reject}})$, where Q, Σ, Γ are **finite** sets and:

1. Q is the set of states,
2. Σ is the input alphabet **not containing the blank symbol** \sqcup ,
3. Γ is the tape alphabet where $\sqcup \in \Gamma$ and $\Sigma \subseteq \Gamma$,
4. $\delta: Q \times \Gamma \rightarrow Q \times \Gamma \times \{L, R, S\}$ is the transition function,
5. $q_0 \in Q$ is the start state,
6. $q_{\text{accept}} \in Q$ is the accept state, and
7. $q_{\text{reject}} \in Q$ is the reject state, where $q_{\text{reject}} \neq q_{\text{accept}}$.

Recognizable/Decidable Languages

M **accepts** $w \in \Sigma^*$ if $\exists C_1, C_2, \dots, C_t$ such that

1. C_1 is the starting configuration of M on w
2. $C_i \rightarrow C_{i+1}$ is a valid step of the TM (for $i = 1, 2, \dots, t - 1$)
3. C_t is an accepting configuration

$$L(M) = \{w \in \Sigma^* : M \text{ accepts } w\}$$

TM M **recognizes** a language $L \subseteq \Sigma^*$ iff for all inputs $w \in \Sigma^*$

1. If $w \in L$ then M accepts w and
2. If $w \notin L$ then M **either rejects w or never halts**

Such languages are called **(Turing)-Recognizable**

TM M **decides** a language $L \subseteq \Sigma^*$ iff for all inputs $w \in \Sigma^*$

1. M **halts** on w
2. M **accepts w iff $w \in L$**

Such languages are called **(Turing)-Decidable**

Church-Turing Thesis



Alan Turing



Alonzo Church

*Intuitive notion
of algorithms*

equals

*Turing machine
algorithms*

Can Turing Machines recognize/decide all languages? NO

Time Complexity

A Decidable Language L

Deciders for L :

Inputs

		M_A	M_B	M_C	M_D	M_E	M_F	M_G	M_H	
ϵ		2	2	5	2	3	4	2	2	...
0		2	5	12	2	3	5	12	5	...
1		20	12	14	13	8	19	2	9	...
00		32	14	18	9	18	3	5	90	...
01		12	21	56	8	12	18	18	30	...
10		21	22	26	15	11	12	32	15	...
11		11	12	25	100	13	48	98	29	...
000		320	201	159	201	190	200	180	65	...
001		211	208	190	200	189	301	219	82	...
010		328	271	214	441	193	208	109	77	...
011		227	261	191	201	188	107	211	207	...
\vdots		\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\ddots

Number of configurations TM needs to reach an accept/reject state on this input

How to compare different deciders?

Two Deciders for L

	M_A				M_B		
ε	2	}	2	ε	2	}	2
0	4	}	12	0	9	}	10
1	12	}		1	10	}	
00	15	}		00	40	}	
01	22	}	110	01	30	}	45
10	17	}		10	45	}	
11	110	}		11	39	}	
000	49	}	300	000	73	}	85
001	34	}		001	77	}	
010	38	}		010	85	}	
011	300	}		011	80	}	
\vdots	\vdots			\vdots	\vdots		



Worst case
in this group

Running Time of a TM

Definition: Let M be a TM that halts on all inputs (**decider**). The **running time** or **time complexity** of M is the function $t: \mathbb{N} \rightarrow \mathbb{N}$ where

$$t(n) = \max_{w \in \Sigma^*; |w|=n} \text{number of steps } M \text{ takes on } w$$

- M runs in time $t(n)$
- n represents the input length

$$t(0) = 2$$

$$t(1) = 10$$

$$t(2) = 45$$


$$t(3) = 85$$

⋮

ε	2	}	2
0	9		
1	10	}	10
00	21		
01	30	}	45
10	45		
11	33		
000	73		
001	77	}	85
010	85		
011	80		
\vdots	\vdots		

Two Deciders for L


M_A



Length 0	}	$t_1(0) = 3$
Length 1		$t_1(1) = 5$
Length 2		$t_1(2) = 9$
Length 3	}	$t_1(3) = 17$
Length 4		$t_1(4) = 33$
Length 5	}	$t_1(5) = 65$
\vdots		\vdots

M_B

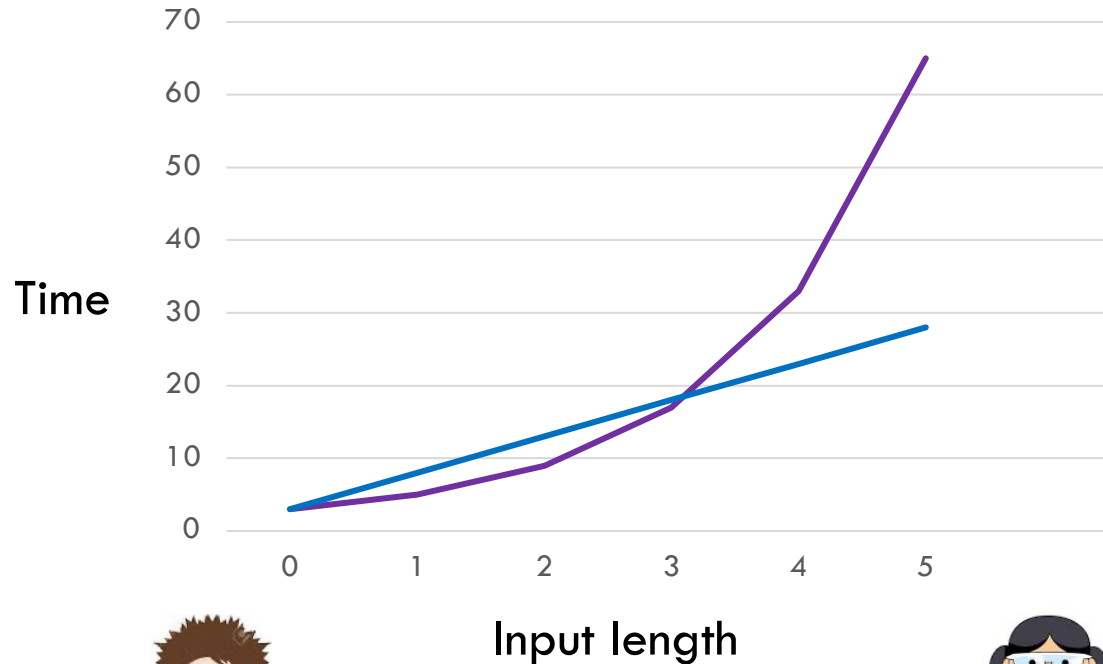
Length 0	}	$t_2(0) = 3$
Length 1		$t_2(1) = 8$
Length 2		$t_2(2) = 13$
Length 3	}	$t_2(3) = 18$
Length 4		$t_2(4) = 23$
Length 5	}	$t_2(5) = 28$
\vdots		\vdots



$$t_1(n) = 2^{n+1} + 1$$

vs

$$t_2(n) = 5n + 3$$



How to compare running time functions?

$$f_1(n) = 2^n, \quad f_2(n) = 5n^3 + 1, \quad f_3(n) = 20n + 6$$

Big-O and Small-o Notation

Definition (Big-O): Let $f, g: \mathbb{N} \rightarrow \mathbb{R}_{\geq 0}$. We say $f(n) = O(g(n))$ if
 $\exists C > 0 \exists n_0$ s.t. $\forall n \geq n_0 \quad f(n) \leq C \cdot g(n)$

Examples:



$$5n^3 + 1 =? O(2^n)$$



$$5n^3 + 1 =? O(20n + 6)$$

Definition (Small-o): Let $f, g: \mathbb{N} \rightarrow \mathbb{R}_{\geq 0}$. We say $f(n) = o(g(n))$ if
 $\forall c > 0 \exists n_0$ s.t. $\forall n \geq n_0 \quad f(n) \leq c \cdot g(n)$

Examples:



$$\sqrt{n} =? o(n)$$




$$f(n) =? o(f(n))$$

$$f_1(n) = 2^n, \quad f_2(n) = 5n^3 + 1, \quad f_3(n) = 20n + 6$$


$$f_3(n) = O(f_2(n)) \quad f_2(n) = O(f_1(n))$$

To Summarize ..

Memory Address	Block Label	Time Step
2	Length 0	$t(0) = 2$
4	Length 1	$t(1) = 4$
3	Length 2	$t(2) = 3$
2	Length 3	$t(3) = 7$
3		
\vdots		
4		
7	Length 4	$t(4) = 15$
\vdots		
6		
10		
15	Length 5	$t(5) = 14$
\vdots		
10		
8		
14		
\vdots		



$$t(n) = O(2^n)$$



Time Complexity

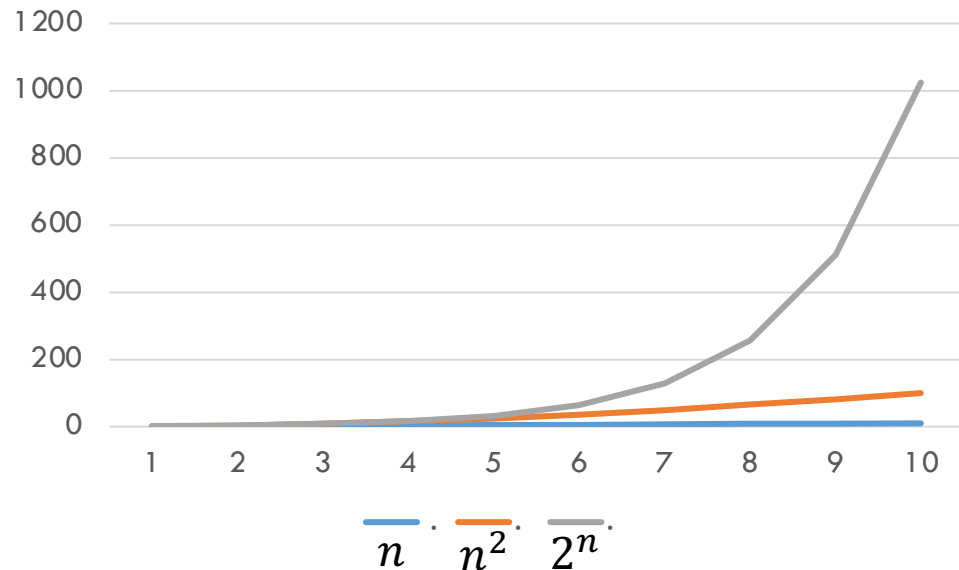
Definition: Time complexity class

$\text{TIME}(t(n)) := \{L \subseteq \Sigma^* \mid L \text{ is decided by a TM with running time } O(t(n))\}$

– $B \in \text{TIME}(n^2)$

Theorem:

$\text{TIME}(n) \subseteq \text{TIME}(n^2) \subseteq \dots \subseteq \text{TIME}(2^{\sqrt{n}}) \subseteq \text{TIME}(2^n) \subseteq \text{TIME}(2^{2^n}) \dots$



The Complexity Class P and Efficiency

Definition: P is the class of languages that are **decidable** in **polynomial time** on a **deterministic single-tape** Turing machine. In other words,

$$P = \bigcup_{k=1}^{\infty} \text{TIME}(n^k).$$

For instance: P is the same class of languages for TMs with 2 tapes.

1. P is invariant for all models of computation that are polynomially equivalent to deterministic single-tape TM – **robust**
2. P roughly corresponds to the class of problems that are realistically solvable – **and we focus on such problems in the course**